

Body-Facing Experiment Tables

Anonymous Author(s)

Submission Id: INTERNAL PRE-RESULTS – NOT FOR SUBMISSION

Abstract

This review copy fixes the compact, body-facing result-table layout. It keeps method rows and primary endpoints visible while moving roots, seeds, budgets, scorers, and execution metadata to the experiment manifest. Dashes reserve empirical cells; they are not zeros or predicted values.

CCS Concepts

• **Computing methodologies** → **Multi-agent systems**; *Cooperation and coordination*.

Keywords

multi-agent systems, role formation, peer judgment, online adaptation, experiments

ACM Reference Format:

Anonymous Author(s). 2027. Body-Facing Experiment Tables. In *Proc. of the 26th International Conference on Autonomous Agents and Multiagent Systems (AAMAS 2027), Hanoi, Vietnam, 3–7 May 2027*, IFAAMAS, 2 pages.

The detailed nine-cell execution contract remains in EXPERIMENT_MATRIX_EXECUT... The body tables below show only the comparisons and endpoints that a reviewer needs to read the claims. Each numerical entry will be reported as a stream-level estimate with its pre-registered interval; a dash is an unfilled cell.

- 1 **Primary results**
- 2 **Online service and stability**
- 3 **Mechanism and responsibility checks**



This work is licensed under a Creative Commons Attribution International 4.0 License.

Proc. of the 26th International Conference on Autonomous Agents and Multiagent Systems (AAMAS 2027), M. Baldoni, F. Fang, W. Yeoh, N. Yorke-Smith (eds.), 3–7 May 2027, Hanoi, Vietnam. © 2027 International Foundation for Autonomous Agents and Multiagent Systems (www.ifaamas.org).

Policy	RQ1: situated signal		RQ2: future assignment		Cost	Streams
	Brier ↓	Calibration ↓	Quality–cost utility ↑	Assignment change ↑	Complete cost ↓	
Uniform	—	—	—	—	—	—
No update	—	—	—	—	—	—
Raw acceptance	—	—	—	—	—	—
Terminal only	—	—	—	—	—	—
Ctx.-trust linear	—	—	—	—	—	—
Pooled controller	—	—	—	—	—	—
RARE	—	—	—	—	—	—

Table 1: Main matched-information comparison for situated judgment and future assignment (RQ1–RQ2). Entries are mean $\pm$ the interval fixed in the confirmation card; best values will be bold only after the complete-cost and independent-outcome checks pass.

Policy	Payoff ↑	Regret ↓	Update p95 ms ↓	State KB ↓	Recovery steps ↓	Streams
Uniform	—	—	—	—	—	—
No update	—	—	—	—	—	—
History only	—	—	—	—	—	—
Ctx.-trust linear	—	—	—	—	—	—
RARE	—	—	—	—	—	—

Table 2: PeerSelect service and drift comparison (RQ3). Arrival rate, selected-only payoff, model route, and resource budget are fixed across rows.

Mechanism ablations				
Variant	Future utility ↑	Assignment change ↑	False attribution ↓	Complete cost ↓
Full RARE	—	—	—	—
No public evidence	—	—	—	—
No delayed credit	—	—	—	—
No ownership gate	—	—	—	—
No selected-only rule	—	—	—	—
No correction / replay	—	—	—	—
Responsibility and transport stress				
Scenario	False attribution ↓	UNKNOWN precision ↑	No-op rate ↑	Recovery steps ↓
Recipient-only mutation	—	—	—	—
Mixed ownership mutation	—	—	—	—
Delayed / out-of-order	—	—	—	—
Resource failure	—	—	—	—

Table 3: Mechanism and responsibility evidence (RQ4). The upper panel tests which components change future assignment; the lower panel tests whether ownership, delay, and resource failures fail closed rather than creating producer credit.