

Know Who You Are: Earning Roles from Situated Peer Judgments

Anonymous Author(s)

Submission Id: INTERNAL PRE-RESULTS – NOT FOR SUBMISSION

Abstract

Multi-agent workflows often learn who should receive the next responsibility from acceptance signals or terminal success. Those signals are unsafe when a recipient must integrate a producer’s artifact: the same failure can arise from a producer defect, recipient-owned work, or a downstream component. We identify this failure as *responsibility confounding* and introduce *earning roles*, a protocol that turns a recipient’s situated use into version-bound, responsibility-scoped evidence. A gate separates producer-attributable evidence from recipient work and unresolved cases; a later selector reads only the eligible snapshot, seals its assignment before execution, and applies one delayed selected-only update after the target outcome arrives. This ordering yields testable noninterference, attribution-safety, and replay invariants. We define an ArtifactRole track for producer–recipient evidence and a complementary PeerSelect track for local selection and drift, with matched information, candidate menus, feedback delays, and complete-cost accounting. The resulting evaluation asks whether situated evidence improves future assignment quality beyond acceptance, terminal-only, and same-information contextual baselines, while measuring latency, forgetting, and false attribution. The scope is responsibility adaptation in dependency-bearing workflows, rather than universal expertise or autonomous workflow discovery.

CCS Concepts

• **Computing methodologies** → **Multi-agent systems**; *Cooperation and coordination*.

Keywords

multi-agent systems, role formation, peer judgment, online adaptation, responsibility

ACM Reference Format:

Anonymous Author(s). 2027. Know Who You Are: Earning Roles from Situated Peer Judgments. In *Proc. of the 26th International Conference on Autonomous Agents and Multiagent Systems (AAMAS 2027)*, Hanoi, Vietnam, 3–7 May 2027, IFAAMAS, 10 pages.

1 Introduction

In a dependency-bearing workflow, a collaborator’s value is observed through another agent’s attempt to use its artifact. Consider two queue-processing cases with the same rejection and failed team outcome. In one, a producer violates the priority-module contract; in the other, the module is correct but the recipient implements its consumer incorrectly. A terminal score or a bare rejection therefore points to the wrong future responsibility in at least one case. The

recipient is closest to the delivery, yet the same integration step can expose a producer defect or create recipient-owned work.

This is *responsibility confounding*: one feedback channel mixes producer correctness, recipient integration, and downstream outcome. Reputation and partner-selection systems make adaptive choice possible, but their usual interaction-level signal does not identify which side of a concrete producer–recipient dependency caused the failure [10, 13–15]. The problem is not solved by adding a larger score. The evidence must say what was observed, whose paths changed, and when that evidence became legal to use.

We ask: *when can the agent that actually receives and uses a delivery provide reliable evidence about the producer’s future responsibility?* The answer has a necessary sequence:

delivery → recipient use and judgment → producer attribution
→ future assignment → unseen quality and cost.

(1)

A final score cannot identify which producer caused a defect. An acceptance bit cannot distinguish a useful artifact from an easy integration. A recipient’s normal integration work cannot become a producer penalty without an ownership check. Conversely, a role posterior that never changes a later pre-execution assignment is not role formation.

Our central operation is to make the recipient’s local observation a public, version-bound evidence object with a responsibility boundary. A producer is credited only when the delivery was observed, the recipient action is traceable to that delivery, and a contract check excludes recipient-owned work. Publication does not silently update the policy. Before a later task starts, a legal owner reads the snapshot and the selector seals its assignment and propensity. Only after the later task has an independent outcome can one delayed credit update occur. The protocol therefore separates what a recipient saw from what a future responsibility produced, while keeping the separation visible in a replayable ledger.

Figure 1 shows the loop; its colors separate situated observation, public evidence, and the later decision.

We call this protocol *earning roles*: a collaborator earns a future responsibility opportunity through attributable evidence that changes assignment under a fixed information and cost contract. We use *responsibility-aware role evidence* (RARE) for the protocol’s evidence semantics, not for a particular encoder, head, RLS update, or periodic refit. The primary claim is deliberately narrow: situated evidence is useful only if it adds producer-attributable information beyond raw acceptance, terminal-only outcomes, and an equally informed contextual selector.

The paper makes four contributions.

- We formulate responsibility confounding as an information and attribution problem, separating producer quality, recipient action, downstream adoption, final outcome, and UNKNOWN while preserving delivery, version, read cut, and propensity.



This work is licensed under a Creative Commons Attribution International 4.0 License.

- We introduce the earning-roles protocol and three audit invariants: target outcomes cannot affect a sealed selection, recipient-only work cannot create producer evidence, and replay cannot credit one assignment twice.
- We define an ArtifactRole track and a secondary PeerSelect mechanism track, with root-level splits, same-information baselines, and complete resource accounting that distinguish evidence value from closed-loop team benefit.
- We make the claim falsifiable: no out-of-sample information, an equally informed contextual explanation, unsafe attribution, or a gain that disappears at matched cost invalidates the role-formation interpretation.

The remainder of the paper follows this causal chain. Related work identifies the missing responsibility semantics; the next sections define the event, gate, and update order; the benchmark section fixes the candidate tracks and controls; and the experiment section maps each mechanism claim to an estimand, a result cell, and a falsification rule.

2 Related Work and Positioning

The nearest work differs from our question along three axes: what is observed, when state changes, and whose contribution receives credit. Table 1 makes that boundary explicit. Partner-selection and reputation methods motivate adaptive choice, while graph and role systems motivate changing organization; neither, by itself, identifies whether a concrete producer delivery caused a recipient’s action. Recent collaboration benchmarks measure team-level coordination and competition, while interactive agent benchmarks measure environment-level decisions and tool/user outcomes; neither by default turns a recipient’s integration trace into a producer-owned label [1, 9, 18, 21]. A contextual-trust or delayed-bandit adapter supplies a strong same-information alternative, not a synonym for the proposed mechanism.

The distinction from ordinary online learning is temporal and semantic. LLM collaboration frameworks make multi-agent conversation and workflow control practical, but their task-level success signal still leaves this ownership question open [16, 22]. A source judgment is available before a later assignment, but the later quality outcome is available only after that assignment has executed. A target outcome may update later decisions once it arrives. Using it to reconstruct an earlier assignment would violate that assignment’s information boundary; treating it as a label of the source artifact would also require a separate attribution argument. Our protocol makes the two events explicit and reports the service cost of maintaining the separation. Bandit estimators therefore provide useful update comparators; they are not themselves the claimed contribution.

Cooperative multi-agent learning has long treated credit assignment as a separate problem, using counterfactual advantages or value factorisation to distribute a team payoff [3, 11]. Those methods establish why a shared outcome is insufficient for individual credit, but their reward/state contracts do not identify ownership of a concrete artifact inside a recipient’s integration path; this is the boundary tested here.

3 Problem Formulation

3.1 Objects and information boundary

At episode t , the task supplies context x_t and a legal candidate set C_t . The selector chooses one candidate before any candidate in C_t executes:

$$a_t \sim \pi_\theta(\cdot \mid x_t, C_t, R_t), \quad a_t \in C_t, \quad (2)$$

where R_t is the public evidence snapshot available at read cut w_t . Only the selected candidate is charged for execution. The candidate emits a versioned delivery o_t with digest d_t ; a recipient j_t reads that delivery and takes action m_t . The recipient may accept, revise, reject, or be unable to complete the integration.

The smallest auditable event is

$$e_t = (x_t, C_t, a_t, o_t, j_t, J_t, m_t, Q_t, Y_t, \tau_t, c_t), \quad (3)$$

where J_t is the recipient judgment, Q_t is the producer-contract check, Y_t is an independent later or terminal outcome when the task defines one, τ_t is feedback-arrival metadata, and c_t is the complete cost record. The event also stores the candidate version, task/root identifier, artifact digest, feature-schema version, selection probability $p_t = \pi_\theta(a_t \mid x_t, C_t, R_t)$, read-cut digest, and evidence lineage.

Concrete instance. In the queue case, x_t is the JSON task request, $C_t = \{A, B\}$ is the legal producer menu, and $a_t = B$ selects producer B . The delivery o_t is B ’s versioned priority-module artifact with digest d_t ; recipient field $j_t = k$ records k ’s import, test, or repair action as m_t ; Q_t checks only the producer-owned queue and priority paths; Y_t is the later consumer outcome; and c_t contains model, tool, token, and wall-clock cost.

A repair to k ’s consumer path is recorded in m_t and sets the responsibility gate to $G_t = 0$ rather than changing B ’s producer evidence.

The policy may consume only public events whose arrival is at or before its read cut. Hidden gold, a private scorer, a future target outcome, another policy’s state, and an unexecuted candidate cannot affect selection, retry, stopping, or update. A missing or non-attributable event is UNKNOWN, not a negative label. This boundary is the difference between a role signal and a convenient retrospective score.

3.2 Responsibility gate

For a source episode, define $G_t = 1$ only when all of the following hold: (i) the frozen contract, candidate registry, changed paths, and independent producer check yield a unique producer structural owner (or a pre-registered source defect is uniquely bound to the producer); (ii) the recipient inspected o_t and its action is traceable to o_t ; (iii) recipient-owned integration paths are empty; (iv) J_t , m_t , Q_t , identifiers, versions, and digests are complete; and (v) the source event has not already been credited. The judge’s free-text target role is retained as a calibration observation and its disagreement with the structural owner is logged, but it cannot by itself create or cancel evidence. Otherwise $G_t = 0$ and the source cannot create producer evidence.

This gate distinguishes four observations that are often collapsed: a producer-content defect, a recipient’s own integration work, a sink or final-task failure, and an unresolved attribution. The public




Figure 1: The earning-roles protocol. Blue marks peer context and the read cut; orange marks situated observation; green marks producer-attributable evidence and the later outcome; violet marks the future assignment. The UNKNOWN branch remains audit-only. Gray dashed feedback is delayed and selected-only, entering a future state and next read cut without changing historical evidence or a sealed assignment. Peer-context links are descriptive, not causal.

Representative method and problem	Observed signal and update	Responsibility / credit semantics	Experimental object	Difference / test here
Curiosity-Driven Partner Selection; Defection at First Sight [13, 15]	Interaction outcomes; partner state changes after social-game feedback	Peer-level interaction payoff; the cited formulations do not expose a producer-recipient artifact gate	Convention or dilemma games	Versioned delivery use and producer attribution
Reputation as a Solution to Cooperation Collapse [12]	Shared assessments and cooperation history; pooled reputation update	Broad reputation/cooperation score; the cited formulation does not expose producer and recipient ownership fields	Language-agent social cooperation	Public responsibility-safe evidence and explicit UNKNOWN
AgentNet; GPTSwarm [17, 22]	Agent traces, graph structure, or routing outcomes; organization changes over time	Route/role utility; the cited formulations do not define a producer-artifact credit label	Graph or workflow coordination	Recipient action bound to one delivery and future-owner assignment
Dynamic Role Assignment; MorSE [8, 20]	Task context and role/subtask performance; roles or experts are reassigned	Role/expert utility; the cited formulations do not expose producer-artifact attribution inside recipient integration	Debate or mixture-of-experts workflows	Responsibility gate, cross-owner propagation, and delayed correction
Finite-time multi-armed bandit [2]	Selected-arm rewards; update after an observed arm payoff	Arm-level reward; no producer-recipient ownership contract	Pre-execution selection under observed payoff	Classical selected-reward control; contextual and delayed adapters require separate qualification

Table 1: Novelty boundary used in the paper. The rows are comparison families, not claims that each source implements the same task. A faithful adapter must preserve the listed visible information, arrival order, selected-only denominator, and complete cost before entering the scientific matrix.

record stores the reason for $G_t = 0$. This makes safety measurable: a method that refuses ambiguous evidence may have a higher UNKNOWN rate while producing fewer false producer penalties.

Figure 2 shows the episode, public-evidence, and local-decision lanes.

4 Earning Roles: A Responsibility-Aware Protocol

The equations below specify the candidate realization used to qualify the protocol interfaces; they do not lock the final backbone or updater. Those choices remain an experimental question and will be fixed only after the same-information controls expose the relevant bottleneck.

4.1 Two-stage evidence and update

The protocol has two independent interfaces:

$$\text{publish}(e_t) \rightarrow \{\text{RoleEvidence}, \text{UNKNOWN}, \text{invalid}\}, \quad (4)$$

$$\text{choose}(x_t, C_t, R_t) \rightarrow (a_t, p_t, \delta_t), \quad (5)$$

$$\text{validate}(j_t) \rightarrow \{\text{credit}, \text{UNKNOWN}, \text{invalid}\}, \quad (6)$$

$$\text{update}(\text{credit}) \rightarrow \{\text{updated once}, \text{no-op}\}. \quad (7)$$

Publication may change the public input to a later assignment, but it does not call the persistent updater. For a later target assignment j , credit is legal only if the assignment was written before target selection and task start, the selected candidate and propensity match the sealed preview, the target delivery and outcome are complete, and the assignment has not already been credited. The update then evaluates that future responsibility opportunity, never the source episode.




Figure 2: Evidence-to-role state. The episode lane records what the producer and recipient actually do; the public-evidence lane projects only typed, responsibility-checked fields; and the local-decision lane reads eligible rows at a read cut before sealing an assignment. Target outcomes arrive after the seal and update only the matching selected opportunity. The representation and updater remain replaceable interfaces.

The event-time contract is therefore

$$\begin{aligned} &\text{source selection} \rightarrow \text{delivery} \rightarrow (Q, J, m, Y) \rightarrow G \\ &\rightarrow \text{evidence publication} \rightarrow \text{assignment} \rightarrow \text{target outcome} \quad (8) \\ &\rightarrow \text{delayed selected-only update.} \end{aligned}$$

A correction creates a superseding evidence version and affects only future snapshots. Duplicate, late, out-of-order, resource-failed, or invalid events are idempotent no-ops or explicit UNKNOWN records. Ledger replay must reproduce the same state digest and choice sequence.

4.2 Evidence representation and candidate score

A minimal public state for producer v in context r is

$$R(v, r) = (n(v, r), s(v, r), u(v, r)), \quad (9)$$

where n counts eligible evidence, s accumulates responsibility-safe positive/negative labels, and u records unresolved or rejected evidence for audit. A candidate score can combine a task representation $h(x, v, r)$ with the public evidence term:

$$\text{score}(v \mid x, r) = w^\top h(x, v, r) + \beta \frac{s(v, r)}{n(v, r) + \alpha}. \quad (10)$$

For the queue instance, r is the priority-contract context, $h(x, B, r)$ is the fixed task-candidate representation for producer B , and n, s, u are counts over the versioned events that passed the gate. A

recipient-only repair increases u or creates an explicit UNKNOWN record; it cannot increase s for B .

The representation h , updater, and constants are replaceable implementation conditions rather than the scientific object. The score is a contract-level construction, not a theorem or a claim that one backbone is optimal. The selector uses a pre-registered exploration floor, records p_t , and uses preview-assignment-commit so that writing a future assignment cannot resample the selected candidate.

The same protocol yields three observable predictions. First, responsibility safety is measured by false producer-attribution and UNKNOWN rates under recipient-only interventions. Second, temporal validity is measured by read-cut and assignment-order invariants and by the absence of target-outcome fields at selection. Third, bounded online service is measured by per-event update latency, backlog, state bytes, and complete cost. These predictions map to the RQ2–RQ4 endpoints; they are not assumptions that the score automatically satisfies.

4.3 Executable selector and update contract

The protocol becomes runnable with a small stateful selector whose interfaces are fixed before confirmation data are read. We use a frozen text encoder only to place a task, a candidate’s public description, and its eligible evidence in a common space; the encoder is

Table 2: State fields and their legal update event. A field is read only after its listed cut and never retroactively edits a sealed assignment.

Field	Source	Update boundary
Profile $P_{v,k}$	Gate-passed evidence before κ_k	Next profile cut; immutable thereafter
Delta $D_{v,k,t}$	New gate-passed evidence after κ_k	evi- Read watermark w_t
Residual θ_v	Selected target outcome Y_t	Matching delayed update only
Uncertainty $u(v)$	Unknown, duplicate, or correction records	Monotone audit update
Assignment a_t, p_t	Candidate menu and snapshot	Sealed before target execution

not updated by a feedback event. To prevent a profile from silently reusing the label that trains the decision head, each candidate has a timestamped profile snapshot and a post-snapshot evidence delta:

$$P_{v,k} = \text{summarize}\{e_i : G_i = 1, t_i < \kappa_k\}, \quad (11)$$

$$D_{v,k,t} = \text{summarize}\{e_i : G_i = 1, \kappa_k \leq t_i \leq w_t\}, \quad (12)$$

$$z_t(v) = [\phi(x_t) \parallel \phi(P_{v,k}) \parallel \phi(D_{v,k,t}) \parallel b(v)]. \quad (13)$$

Here κ_k is the last profile cut, w_t is the legal read cut, ϕ is frozen, and $b(v)$ contains version and uncertainty features that are already public at w_t . An empty profile or delta is represented explicitly; it is not replaced by a future label. The split is a provenance rule, not a tuning convention: an event can enter exactly one of P , D , or a later target outcome.

The candidate distribution is

$$q_t(v) = \text{softmax}_{v \in C_t} (g_\psi(z_t(v)) + \theta_v^\top \phi(x_t)), \quad a_t \sim q_t, \quad (14)$$

where g_ψ is a fixed small head and θ_v is a per-candidate residual that is allowed to change only after an independent outcome for a previously selected assignment arrives. The head, temperature, exploration floor, clipping rule, and learning rate are configuration fields, not values selected after seeing confirmation results. The experiment compares the full state against profile-only, delta-only, residual-only, and no-history versions, so a gain cannot be attributed to the word “fusion” alone.

At runtime the selector executes five auditable operations. (1) It appends a source judgment to an immutable ledger and applies the ownership gate; rejected or incomplete records become UNKNOWN. (2) It advances a read watermark and materializes a state snapshot, recording the feature-schema and state digest. (3) It scores only the legal candidate menu and writes the candidate set, probabilities, sampled action, and assignment digest before the target starts. (4) When the target outcome arrives, it verifies the matching assignment, selected-only eligibility, version, and idempotency key. (5) It applies one bounded residual update and emits a new state digest; profile and delta materialization wait for the next declared read cut. This sequence permits a crashed service to replay events without changing a sealed choice.

This implementation contract gives the method a measurable online meaning without claiming that one encoder or head is universally best. It also makes the difficult cases executable: a late judgment changes a future snapshot, a recipient-owned repair cannot train the producer residual, and a missing target outcome produces no update. The confirmation experiment fills in latency, state size,

and predictive utility from the same ledger used to audit attribution. If a simpler calibrated comparator matches the full selector under the same read cuts and complete cost, the paper reports the protocol boundary and does not attribute the gain to the fusion architecture.

5 Benchmark and Baselines

5.1 Two complementary tracks

The primary track, **ArtifactRole**, is a TeamBench-derived candidate benchmark [5]. It contains producer artifacts, recipient dependencies, use/rework/rejection actions, responsibility checks, later assignments, downstream adoption, and complete cost. Candidate roots are DIST1 queue-race (development diagnostic) and PIPE3 stream-processing (conditional confirmation); renamed seeds are not independent roots. Promotion requires material visibility, write permissions, scorer coverage, root independence, and clean replay. This stricter contract is intentional: TeamBench’s role separation supplies the dependency scaffold, while ArtifactRole adds the producer–recipient attribution and future-assignment fields needed by our question.

Producer conformance and terminal usability are scored separately: a recipient may normalize a nonconforming delivery into a usable result. Every executed arm, including no-update, receives terminal evaluation. A prior evidence-based assignment is required for assignment credit, not for inclusion in the outcome denominator; native benchmark conformance remains a separate endpoint.

The secondary **PeerSelect** track uses a pinned local graph interaction benchmark for selected-only payoff, local candidate sets, online updates, drift, and service cost. It isolates update behavior but cannot establish situated artifact judgment or producer attribution. The tracks share event and selector interfaces while keeping tasks, labels, scorers, denominators, and claims separate.

5.2 Predeclared baselines

All primary arms receive the same candidate menu, initial opportunities, model/API/tool budget, exploration opportunity, feedback arrival schedule, failure handling, and complete cost accounting. Each arm reads only its legal history. Table 3 names the alternative explanations that the matrix must rule out.

If the closest published method cannot preserve the same public fields, candidate menu, arrival order, selected-only rule, and complete-cost contract, the confirmation card records an explicit NO-GO; it does not replace that missing adapter with a self-authored weak control.

RLS, online logistic/SGD, periodic refitting, and any candidate incremental updater are training comparators. They are not counted as separate selection policies and are not presented as the innovation merely because they are fast. The contextual-trust arm follows the same information contract as classical contextual bandits [7]; its delay handling is compared with the general delayed-feedback setting [4]. Agent ranking and human–LLM teaming provide additional positioning points, but neither substitutes for the producer–recipient label contract [6, 19].

6 Experimental Questions and Matrix

The matrix below is a pre-specified design scaffold, not an observed result table. A cell enters the confirmation analysis only after its

Arm	Read/update contract	Alternative explanation tested
Uniform	Candidate menu; fixed local-uniform sampling	Random allocation and exploration lower bound
No update	Initial prior and current context; feedback ignored	Static execution is sufficient
Raw acceptance	Recipient accept/reject only	Responsibility filtering adds no value
Terminal only	Independent final outcome only	Situated judgment adds no information
Contextual trust/bandit	Same context, judgment, propensity, delay, and budget as RARE	The effect is ordinary contextual trust or bandit learning
Pooled controller	Permitted public history under the same total capacity	A shared controller explains the gain
Closest published adapter	Faithful public-information adapter, if it passes qualification; otherwise an explicit no-go cell	Existing role/organization method is sufficient
RARE	Responsibility-safe evidence plus delayed selected-only credit	Candidate closed-loop mechanism

Table 3: Primary comparison boundary. A named arm enters the scientific matrix only after an executable adapter, information declaration, cost contract, and independent qualification pass.

root, stream, seed, budget, scorer, cost contract, stopping rule, and manifest digest are frozen before the corresponding calls; otherwise it remains planned or UNKNOWN.

6.1 Research questions and estimands

RQ1 (information value). Does an eligible situated judgment predict an independent producer-contract or later-use result beyond raw acceptance and terminal-only feedback? The primary estimands are out-of-sample log loss, Brier score, AUC when defined, calibration, and incremental predictive value on an independent root. The confirmation card also freezes a field-level feature manifest: a judgment-only view, a producer-check/terminal-only view, and their declared fusion. Artifact, structural owner, Qp, Y, read cut, arrival schedule, and cost remain fixed while legal recipient-judgment mutations test whether the assignment or score actually responds to J.

RQ2 (assignment causality). Does public evidence change an assignment made before target execution, including when the future owner did not produce the source judgment? We measure assignment-change probability, propensity-weighted future quality, adoption, rework, and complete cost. Evidence from producer *B* cannot be consumed as if it belonged to producer *C*.

RQ3 (online behavior). Does the update path satisfy the service contract under selected-only delayed feedback? We report publish, read/assignment, and update p50/p95 latency, service lag, backlog, state bytes, CPU/GPU/RAM, token/tool/API calls, and replay time.

RQ4 (stability and boundaries). How does the method behave under task or peer drift, delayed or out-of-order feedback, version replacement, recipient-only work, mixed ownership, scorer failure, and resource failure? We report old-root peak and mean forgetting, drift recovery window, UNKNOWN rate, correction consistency, and false producer attribution.

The primary unit of uncertainty is an independently reset stream. Episodes inside one stream are dependent. We use paired stream-level estimates, bootstrap or randomization intervals over streams, and pre-registered multiplicity control for the primary contrasts. Missing, invalid, or unknown events remain in the denominator and are never silently recoded as failures or successes.

Primary endpoint design. The design reserves four card identifiers: H1 uses independent-target Brier score (lower is better) for incremental situated signal; H2 uses future assignment quality-complete-cost utility (higher is better) against the same-information control; H3 uses selected-feedback update p95 latency and backlog under the fixed arrival budget (lower is better); and H4 uses false producer-attribution rate under ownership mutations (lower is better). Each card is required to fix its control, minimum practically important difference or precision target, exploration floor/positivity condition, stream-level 95% interval, and stopping rule before outcomes are observed; the current design-only manifest records which fields remain unfrozen.

Only these H1–H4 endpoints are primary for the headline interpretation. The remaining measures in the matrix are secondary or exploratory and cannot be promoted after outcomes are observed.

6.2 Complete experiment matrix

Table 4 is the execution skeleton. The result columns are deliberately empty; the confirmation manifest must bind every row to a root, split, policy, seed, budget, scorer, and raw log before any result is filled.

Figure 3 links the two tracks to the four research-question endpoints; it does not merge their labels or claims.

Table 5 reserves the endpoints that will carry the paper’s headline comparison once the confirmation manifest is complete.

6.3 Budget, statistics, and reporting

Every model request reserves attempted calls, prompt tokens, maximum completion tokens, tool time, and storage before dispatch; errors release only unused reservations and retain their observed cost. We charge producer execution, recipient integration, judgment service, communication, retries, replay, verification, and returned artifacts. Development/search cost is reported separately from deployment cost. A pooled controller receives the same total capacity, and arms that do not use a judgment still pay the declared comparison cost or use a shared judgment call whose output is masked.

For each stream we record configuration, model route and version metadata, code and data hashes, seed, hardware, timestamps, raw API responses, token usage, intermediate actions, failure traces, UNKNOWN reasons, and per-sample outcomes. We report means

RQ/cell	Track/root	Policy/intervention	Primary comparison	Measures	Result
RQ1-I	ArtifactRole / dev root	raw acceptance vs terminal only vs RARE	Same producer deliveries and recipient observations; legal feedback only	Brier/log loss, calibration, judge reliability	—
RQ1-II	ArtifactRole / confirmation root	contextual trust vs RARE	Same public read cut, delay, propensity, model and budget	Incremental prediction and cross-root transfer	—
RQ2-I	ArtifactRole / confirmation root	no evidence vs public evidence only	Preview-assignment before target execution	Assignment change, future quality, adoption, rework	—
RQ2-II	ArtifactRole / confirmation root	delayed update only vs public evidence + delayed update	Same target tasks and selected-only outcome	Future utility and complete cost	—
RQ2-III	ArtifactRole / confirmation root	RARE vs pooled controller	Same total public capacity and candidate menus	Quality-cost utility, concentration, propensity	—
RQ3-I	PeerSelect / dev streams	uniform, no-history, history-only, candidate updater	Local graph, selected-only payoff, fixed call/token budget	Regret/payoff, update p50/p95, state bytes	—
RQ3-II	PeerSelect / drift streams	updater with/without correction and replay	Pre-registered peer/task drift and delayed feedback	Recovery window, forgetting, backlog, replay equality	—
RQ4-I	ArtifactRole / all qualified roots	responsibility mutation matrix	Producer-only, recipient-only, mixed, unknown, resource failure	False attribution, UNKNOWN precision, no-op rate	—
RQ4-II	ArtifactRole / all qualified roots	leave-one-mechanism-out	Remove gate, public evidence, delay, locality, or updater	Main utility, safety, cost, interaction effect	—

Table 4: Pre-specified experiment matrix. Each row binds a research question to a root, split, policy, seed, budget, scorer, cost record, and raw event manifest before interpretation; a row is not an observed result until its confirmation card is frozen and executed.




Figure 3: Experiment map. ArtifactRole measures the situated producer-recipient phenomenon; PeerSelect isolates local selection, drift, and service cost. All policy arms receive the same candidate menus, lawful feedback schedule, task opportunities, and complete-cost budget before the four research questions are scored.

and medians with 95% stream-level intervals, effect sizes for pre-registered contrasts, per-root heterogeneity, ITT and per-protocol denominators, and quality-cost Pareto points. The main endpoint is assignment-level future quality-complete-cost utility; source producer score is auxiliary and cannot replace it.

For every primary contrast, the experiment card fixes the direction, minimum practically important difference or precision target,

stream-level 95% interval, and stopping rule before confirmation. A row without those fields is a design specification rather than an interpretable result.

6.4 Ablations and falsification

The minimum ablation grid is: (A) no evidence/no update; (B) public evidence only, with persistent state unchanged; (C) delayed update

Table 5: Headline endpoints. Values are supplied only by the frozen confirmation manifest; intervals are stream-level 95% intervals.

Endpoint	RARE	Best same- information arm	Difference / interval
Future quality	—	—	—
Complete cost	—	—	—
Quality–cost utility	—	—	—
Assignment change	—	—	—
Judgment Brier score	—	—	—
Update p95 / state bytes	—	—	—
Old-root forgetting	—	—	—
UNKNOWN / false attribution	—	—	—

only, without new public evidence; and (D) public evidence plus delayed update. Additional interventions remove responsibility filtering, selected-only restriction, delay, locality, correction/replay, or the assignment binding one at a time while holding information, model, candidate menu, and budget fixed. A content mutation that changes a visible judgment while holding terminal quality fixed must change the comparator score if the comparator actually consumes that judgment.

The main claim is withdrawn if an evidence publication implicitly changes persistent policy; if a recipient repair is credited to a producer; if assignment or read-cut order is invalid; if unknown or duplicate events update state; if contextual trust explains the complete effect; or if gains disappear at matched complete cost. A null, a trade-off, or a qualified protocol result is scientifically valid, but it cannot be written as successful role formation.

Phase rule. Before selection, the policy sees only the candidate menu, public-row watermark, read-cut digest, and sealed propensity; hidden gold, unexecuted-candidate outcomes, and future target outcomes remain unavailable. During responsibility assessment, the record contains the judgment, action, owner, changed paths, ownership gate, and explicit UNKNOWN reason. After the assignment seal, only the matching target quality/cost event can produce one selected-only update; a late or duplicate event without correction lineage is not a new label.

6.5 Reproducibility and reporting

Before execution, each experiment card fixes the matrix cell, root, stream, model, budget, scorer, estimand, and stopping rule. A hash-linked ledger connects every aggregate to its raw calls and per-episode outcomes. Failed attempts retain their observed cost; incomplete cells remain TBD, never zero. We report safety without utility gain as a safety result, and prediction without changed future assignment as an information result. Neither alone establishes role formation.

7 Results and Evidence Boundary

RQ1–RQ2 (information and assignment). The confirmation card will fill the held-out Brier score, matched-information comparison, assignment-change rate, and future quality–cost utility in Tables 5 and 4: TBD. Prediction without a sealed future assignment supports an information claim only; assignment without an independent target outcome does not establish responsibility transfer.

RQ3–RQ4 (service and safety). The latency, backlog, state-size, forgetting, UNKNOWN, and false-attribution cells remain TBD. We will interpret a safe null, a quality–cost trade-off, and a qualified positive result as different claims rather than as interchangeable evidence.

Current development traces expose attribution and measurement hazards, but do not establish predictive utility, future assignment gains, or real-time learning. No confirmation cell is filled from these traces. Results require qualified roots, independent streams, matched baselines, and complete cost.

8 Discussion and Limitations

The mechanism addresses responsibility opportunities in tasks with an actual producer–recipient dependency. It does not establish general expertise or autonomous workflow discovery. PeerSelect can isolate local adaptation, but cannot replace ArtifactRole’s attribution evidence; conversely, an ArtifactRole result cannot be interpreted as general bandit performance.

Judge error, selective exposure, root dependence, incomplete responsibility coverage, and judgment cost threaten validity. We therefore preserve UNKNOWN events, log selection propensities, assess judge reliability, and evaluate independent roots and streams. Terminal usability also differs from producer conformance: successful downstream normalization need not vindicate an upstream contract violation. Neither metric alone identifies the producer’s contribution to future team utility.

Backbone and updater selection remains an empirical question. Frozen representations with incremental heads, online logistic/RLS updates, and periodic refitting are matched candidates. Latency alone cannot establish useful adaptation: a fast update may reinforce a misattributed label or forget previously useful distinctions. The drift and old-task evaluations must accompany update timing before a real-time learning claim is supported.

Deployment must expose refusals to learn, preserve correction lineage, and retain enough provenance to replay decisions after evidence compression. Tool selection can share this event interface, but requires its own execution oracle; peer judgments retain the ownership gate. The release will include versioned code, prompts, scorer contracts, raw-event manifests, and table-generation scripts. Offline replay reconstructs decision order; reproducing API latency and token costs requires fresh calls. These boundaries distinguish an executable protocol from demonstrated improvements in role formation.

9 Conclusion

This paper defines a precise object for self-evolving multi-agent roles: a recipient’s situated judgment becomes responsibility-safe evidence, that evidence changes a future assignment before execution, and the later assignment is evaluated on unseen quality and

complete cost. The protocol makes each arrow auditable and gives close alternatives the same information and budget. The current draft records the method contract, benchmark and baseline matrix, statistical plan, and falsification boundary while leaving empirical cells empty. The decisive result will be whether the complete loop improves future responsibility allocation beyond equally informed alternatives without sacrificing attribution safety, online service, stability, or cost.

References

- [1] Saaket Agashe, Yue Fan, Anthony Reyna, and Xin Eric Wang. 2025. LLM-Coordination: Evaluating and Analyzing Multi-agent Coordination Abilities in Large Language Models. In *Findings of the Association for Computational Linguistics: NAACL 2025*. Association for Computational Linguistics, 8053–8072. doi:10.18653/v1/2025.findings-naacl.448
- [2] Peter Auer, Nicolò Cesa-Bianchi, and Paul Fischer. 2002. Finite-time Analysis of the Multiarmed Bandit Problem. *Machine Learning* 47 (2002), 235–256. doi:10.1023/A:1013689704352
- [3] Jakob Foerster, Gregory Farquhar, Triantafyllos Afouras, Nantas Nardelli, and Shimon Whiteson. 2018. Counterfactual Multi-Agent Policy Gradients. In *Proceedings of the Thirty-Second AAAI Conference on Artificial Intelligence*, Vol. 32. doi:10.1609/aaai.v32i1.11794
- [4] Pooria Joulani, Andras Gyorgy, and Csaba Szepesvari. 2013. Online Learning under Delayed Feedback. In *Proceedings of the 30th International Conference on Machine Learning (Proceedings of Machine Learning Research, Vol. 28)*. PMLR, 107–115. <https://proceedings.mlr.press/v28/joulani13.html>
- [5] Yubin Kim, Chanwoo Park, Taehan Kim, Eugene Park, Samuel Schmidgall, Salman Rahman, Chunjong Park, Cynthia Breazeal, Xin Liu, Hamid Palangi, Hae Won Park, and Daniel McDuff. 2026. TeamBench: Evaluating Agent Coordination under Enforced Role Separation. *arXiv preprint arXiv:2605.07073* (2026). <https://arxiv.org/abs/2605.07073>
- [6] Marc Lanctot, Kate Larson, Michael Kaisers, Quentin Berthet, Ian Gemp, Manfred Diaz, Roberto-Rafael Maura-Rivero, Yoram Bachrach, Anna Koop, and Doina Precup. 2025. Soft Condorcet Optimization for Ranking of General Agents. In *Proceedings of the 24th International Conference on Autonomous Agents and Multiagent Systems*. IFAAMAS. <https://www.ifaamas.org/Proceedings/aamas2025/pdfs/p1253.pdf>
- [7] Lihong Li, Wei Chu, John Langford, and Robert E. Schapire. 2010. A Contextual-Bandit Approach to Personalized News Article Recommendation. In *Proceedings of the 19th International Conference on World Wide Web*. ACM, 661–670. doi:10.1145/1772690.1772758
- [8] Peiwen Li, Shiyang Zhang, Yangtian Zhang, Sizhuang He, David van Dijk, and Rex Ying. 2026. MoRSE: Task-Oriented Multi-Agent System with Mixture of Role-Subtask Experts. *arXiv:2608.09251* <https://arxiv.org/abs/2608.09251>
- [9] Xiao Liu, Hao Yu, Hanchen Zhang, Yifan Xu, Xuanyu Lei, Hanyu Lai, Yu Gu, Hangliang Ding, Kaiwen Men, Kejuan Yang, Shudan Zhang, Xiang Deng, Aohan Zeng, Zhengxiao Du, Chenhui Zhang, Sheng Shen, Tianjun Zhang, Yu Su, Huan Sun, Minlie Huang, Yuxiao Dong, and Jie Tang. 2024. AgentBench: Evaluating LLMs as Agents. In *International Conference on Learning Representations*. https://proceedings.iclr.cc/paper_files/paper/2024/hash/e9df36b21ff4ee211a8b71ee8b7e9f57-Abstract-Conference.html
- [10] Iñigo Pinyol and Jordi Sabater-Mir. 2013. Computational Trust and Reputation Models for Open Multi-Agent Systems: A Review. *Artificial Intelligence Review* 40 (2013), 607–626. doi:10.1007/s10462-011-9277-z
- [11] Tabish Rashid, Mikayel Samvelyan, Christian Schroeder de Witt, Gregory Farquhar, Jakob Foerster, and Shimon Whiteson. 2020. Monotonic Value Function Factorisation for Deep Multi-Agent Reinforcement Learning. *Journal of Machine Learning Research* 21, 178 (2020), 1–51. <https://www.jmlr.org/papers/v21/20-081.html>
- [12] Siyue Ren, Wanli Fu, Xinkun Zou, Chen Shen, Yi Cai, Chen Chu, Zhen Wang, and Shuyue Hu. 2026. Reputation as a Solution to Cooperation Collapse in LLM-based MASs. In *Proceedings of the 25th International Conference on Autonomous Agents and Multiagent Systems*. IFAAMAS. <https://www.ifaamas.org/Proceedings/aamas2026/pdfs/UEHN4980.pdf>
- [13] Benedict Russell, Chin wing Leung, and Paolo Turrini. 2026. Defection at First Sight: Learning Partner Selection in Optional Social Dilemmas without Prior Information. In *Proceedings of the 25th International Conference on Autonomous Agents and Multiagent Systems*. IFAAMAS. <https://www.ifaamas.org/Proceedings/aamas2026/pdfs/IBSZ1473.pdf>
- [14] Jordi Sabater and Carles Sierra. 2002. Reputation and Social Network Analysis in Multi-Agent Systems. *Proceedings of the First International Joint Conference on Autonomous Agents and Multiagent Systems*, 475–482. doi:10.1145/544741.544854
- [15] Chin wing Leung, Paolo Turrini, and Ann Nowé. 2025. Curiosity-Driven Partner Selection Accelerates Convention Emergence in Language Games. In *Proceedings of the 24th International Conference on Autonomous Agents and Multiagent Systems*. IFAAMAS, 1282–1290. <https://www.ifaamas.org/Proceedings/aamas2025/pdfs/p1282.pdf>
- [16] Qingyun Wu, Gagan Bansal, Jieyu Zhang, Yiran Wu, Beibin Li, Erkang Zhu, Li Jiang, Xiaoyun Zhang, Shaokun Zhang, Jiale Liu, Ahmed Hassan Awadallah, Ryan W. White, Doug Burger, and Chi Wang. 2023. AutoGen: Enabling Next-Gen LLM Applications via Multi-Agent Conversation. In *arXiv preprint arXiv:2308.08155*. <https://arxiv.org/abs/2308.08155>
- [17] Yingxuan Yang, Huacan Chai, Shuai Shao, Yuanyi Song, Siyuan Qi, Renting Rui, and Weinan Zhang. 2025. AgentNet: Decentralized Evolutionary Coordination for LLM-based Multi-Agent Systems. In *Advances in Neural Information Processing Systems*, Vol. 38. https://papers.nips.cc/paper_files/paper/2025/hash/9a379c1b05793d1c42dc832269834515-Abstract-Conference.html
- [18] Shunyu Yao, Noah Shinn, Pedram Razavi, and Karthik Narasimhan. 2024. *tau-bench*: A Benchmark for Tool-Agent-User Interaction in Real-World Domains. *arXiv:2406.12045* <https://arxiv.org/abs/2406.12045>
- [19] Mireia Yurrita, Davide Dell’Anna, Pradeep K. Murukannaiah, Catholijn M. Jonker, and Pinar Yolum. 2026. Developing Guidelines for Human-LLM Agent Teams: A Multi-Stakeholder Lens. In *Proceedings of the 25th International Conference on Autonomous Agents and Multiagent Systems*. IFAAMAS. doi:10.65109/JOWO4591
- [20] Miao Zhang, Junsik Kim, Siyuan Xiang, Jian Gao, and Cheng Cao. 2026. Dynamic Role Assignment for Multi-Agent Debate. *arXiv:2601.17152* <https://arxiv.org/abs/2601.17152>
- [21] Kunlun Zhu, Hongyi Du, Zhaochen Hong, Xiaocheng Yang, Shuyi Guo, Zhe Wang, Zhenhailong Wang, Cheng Qian, Xiangru Tang, Heng Ji, and Jiaxuan You. 2025. MultiAgentBench: Evaluating the Collaboration and Competition of LLM Agents. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Association for Computational Linguistics, 8580–8622. doi:10.18653/v1/2025.acl-long.421
- [22] Mingchen Zhuge, Wenyi Wang, Louis Kirsch, Francesco Faccio, Dmitrii Khizbullin, and Jürgen Schmidhuber. 2024. GPTSwarm: Language Agents as Optimizable Graphs. In *Proceedings of the 41st International Conference on Machine Learning (Proceedings of Machine Learning Research, Vol. 235)*. PMLR, 62743–62767. <https://proceedings.mlr.press/v235/zhuge24a.html>