Expected Experiment Tables

Point values are conditional expectations for the specified populations. Brackets are **subjective planning ranges**, not measured confidence intervals or calibrated forecasts. Utility is $U = Y - 0.10C$. Complete cost C uses a common uniform-reference unit; Y is future-task partial quality, not native benchmark pass rate.

Main Table 6 — Expected primary results

Endpoint	RARE	Same-info. baseline	Difference / planning range
Future quality	0.608	0.588	+0.020 [-0.016, +0.056]
Complete cost	1.34	1.22	+0.12 [+0.05, +0.19]
Quality–cost utility	0.474	0.466	+0.008 [-0.035, +0.051]
Assignment change	0.20	0.16	+0.04 [-0.025, +0.105]
Judgment Brier	0.203	0.211	−0.008 [-0.020, +0.004]
Update p95 / state bytes	68 / 36,864	52 / 28,672	+16 / +8,192 latency: [+2, +42] ms
Old-root forgetting	0.025	0.041	−0.016 [-0.040, +0.008]
UNKNOWN / false attribution	0.498 / 0.0066	0.498 / 0.0066	0 / 0 identical shared-gate inputs

EXPECTED, NOT OBSERVED

The baseline is contextual-trust-linear throughout, with the same responsibility gate. Quality, cost and assignment are ArtifactRole endpoints; timing and old-root drift are PeerSelect endpoints. Brier uses a common binary producer-contract cohort. Attribution uses the common mutation mixture on page 2; it is not online outcome missingness. Update p95 is the CPU feedback-processing path in milliseconds with cached features, excluding LLM execution. State bytes are selector state, not total RAM. Assignment change is diagnostic. Same-gate safety is an equality check, not a RARE advantage.

Main Table 7 — Expected evaluation–update factorial

Evaluation	Update	Future utility (planning range)	Cost
Masked	Off	0.430 [0.385, 0.475]	1.02
Visible	Off	0.448 [0.401, 0.495]	1.18
Masked	On	0.455 [0.405, 0.505]	1.15
Visible	On	0.474 [0.421, 0.527]	1.34

The visible/on cell is the same RARE population as the primary utility row. The four quality means are 0.532, 0.566, 0.570 and 0.608, respectively. Observation given updating: $0.474 - 0.455 = 0.019$; updating given observation: $0.474 - 0.448 = 0.026$. The interaction is only $0.474 - 0.448 - 0.455 + 0.430 = 0.001$, so these expectations do not posit a synergy gain. The conservative interaction range implied by the four marginal ranges is $[-0.194, 0.196]$; no narrow paired interval is inferred without covariance.

The favorable-condition utility edge is small and uncertain; RARE has higher ArtifactRole cost and higher PeerSelect latency and state. Sparse support reverses its comparison with pooled control, and high load reverses its payoff comparison with contextual trust (page 2). Stronger controls remain near ties: producer-contract Brier targets are 0.207 for task-conditioned ridge and 0.206 for matched composition, versus 0.203 for RARE.

Experiment matrix: expected cell results

Row IDs match the nine-cell manuscript matrix. Each result below refers to its specified cohort; different task or stress populations are never pooled.

Main Table 4 — Expected experimental matrix

Cell	Track / root	Strategy / intervention	Matched comparison	Measurement / expected result
RQ1-I	ArtifactRole / dev	Raw vs. terminal vs. RARE	Same delivery, recipient observation, legal feedback	Brier: .225 / .220 / .203. Log loss: .637 / .625 / .584. ECE: .073 / .068 / .055.
RQ1-II	ArtifactRole / confirm	Contextual trust vs. RARE	Same read cut, delay, propensity, model, budget	Brier: .211 / .203; held-out-root Brier: .231 / .217. Transfer deterioration: .020 / .014.
RQ2-I	ArtifactRole / confirm	No evidence vs. public evidence	Preview and assignment before target execution	Change: .04 / .12; quality: .532 / .566; adoption: .48 / .54; rework: .30 / .24.
RQ2-II	ArtifactRole / confirm	Delayed-only vs. evidence+delay	Same target and selected-only outcome	Utility: .455 / .474; cost: 1.15 / 1.34. Added-evidence utility: +.019.
RQ2-III	ArtifactRole / confirm	RARE vs. pooled controller	Same capacity and candidate menu	Stable utility: .474 / .472; max share: .38 / .40; min propensity: .08 / .06. Sparse utility: .427 / .450.
RQ3-I	PeerSelect / dev stream	Uniform, no/history, updater	Pinned graph and call/token budget	Normal payoff: .500 / .532 / .565 / .606. RARE update p50/p95: 37/68 ms; state: 36,864 B. High-load payoff RARE/context: .542/.558; p95: 480/290 ms.
RQ3-II	PeerSelect / drift	Update with/without replay	Fixed peer/task drift and delay	Correction / none: old-root drop .025/.048; restricted mean recovery 6.5/9.0 updates; recovery by 20 updates .96/.88; backlog p95 1.4/.9; correction/replay agreement 1.00/.86.
RQ4-I	ArtifactRole / eligible roots	Responsibility mutations	Producer, recipient, mixed, unknown, failure	Known-owner mixture: UNKNOWN .498; coverage .502; FA/all .0066; FA/claimed .01315; UNKNOWN precision .9592. Source-publication no-op 1.00; unknown-owner invalid credits 0.
RQ4-II	ArtifactRole / eligible roots	Remove one mechanism	Gate, evidence, delay, locality, updater	Utility full/gate removed/evidence removed/updater off/locality removed: .474/.461/.455/.448/.470. FA full/gate removed: .0066/.1600. Utility interaction: +.001.

UNREVIEWED — NOT OBSERVED

Safety denominator. The fixed owner-truth mixture weights are .50 producer-only, .20 recipient-only, .10 mixed, .10 sink-only and .10 resource failure. Their claim coverage is .90/.04/.18/.06/.20; all-attempt false attribution is .008/.004/.012/.002/.004. Hence coverage is .502, UNKNOWN is .498 and false attribution is .0066; conditional claimed-case error is .0066/.502 = .01315. Unknown-owner cases are checked separately and do not supply invented ground-truth labels. UNKNOWN precision uses independently required abstention, not hidden owner identity; source-publication no-op is a parameter invariant. Nonzero FA means verdict error, not prohibited source-credit updating. Updater-off equals the visible/off factorial cell. Information/update ablations retain identical safety on this fixed corpus.

Service and sampling. Normal selected payoff is $(r + 0.3)/0.6$ for the pinned $[-0.3, 0.3]$ reward contract, with dropped/timed-out arrivals scored zero in the ITT mean; it is not mutual-cooperation rate. Recovery concerns new-task payoff (95% of a frozen reference for five windows, capped at 20 updates); forgetting uses old-root tasks. The online path uses cached features, three local candidates, four logical CPU cores and a 1 GiB RAM budget. The planning design has three roots and six reset streams per root/arm; roots are not interchangeable independent event replicates. Workload assumptions and numeric rationale are versioned with the JSON companion. These tables set comparison expectations, not criteria for choosing or discarding actual experiment outcomes.